

From BRMS to Stan

Mitzi Morris

2025-03-18

Stan Development Team



BRMS: Bayesian Regression and Multilevelmodeling in Stan



brms

Bayesian regression models using Stan

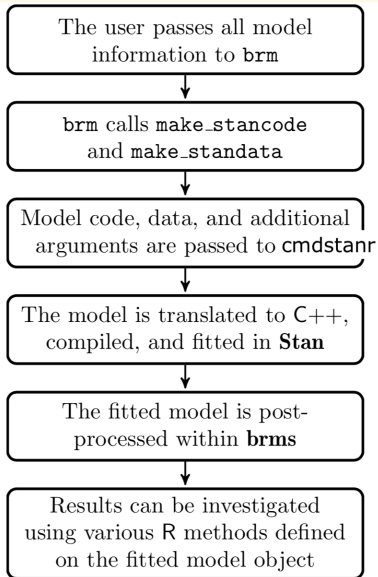
The **brms** package provides an interface to fit Bayesian generalized (non-)linear multivariate multilevel models using Stan. The formula syntax is very similar to that of the package lme4 to provide a familiar and simple interface for performing regression analyses.

<https://paul-buerkner.github.io/brms/>

My Questions for You

- What kinds of models do you use for research? For teaching?
- BRMS
 - How often do you use it: never, sometimes, always?
 - What do you like/dislike/find confusing about BRMS?
- Stan
 - Have you run Stan models through an interface, RStan, CmdStanR, CmdStanPy?
 - Have you written / tried to write or modify Stan models?
- What other tools do you use?

BRMS Processing



Specify and fit a model in BRMS

```
fit1 <- brm(Reaction ~ Days + (Days|Subject),  
           data = sleepstudy)
```

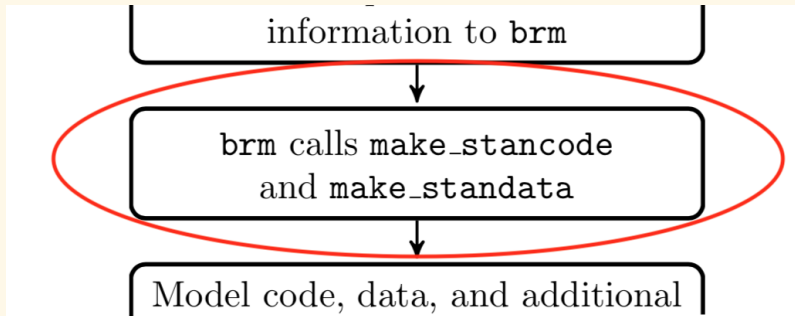
```
sleep_model <- stancode(fit1)
```

```
sleep_data <- standata(fit1)
```

What's the problem?

BRMS Processing

This is the problem!



The Stan program created by `make_stancode` is not easy to understand or modify.

BRMS formula syntax

- The BRMS formula specifies the regression.
- For a multilevel model, BRMS extends the R formula syntax, following lme4.

`response ~ pterms + (gterms | group)`

- The `pterm`s contain *population-level* effects, assumed to be the same across observations.
- The `gterm`s contain *group-level* effects, assumed to vary across the grouping variables specified as `group`.
- The intercept term is either 1 or 0 for no intercept; if unspecified, default is 1.

Function brm

```
fit1 <- brm(Reaction ~ Days + (Days|Subject),  
            data = sleepstudy)
```

- First argument - required: the extended-syntax R formula for the linear predictor.
- Second argument (data) - required: a dataframe containing columns names which correspond to the terms in the formula.
- Additional arguments to the brm function allow further model specification and configure the call to the inference algorithm.
 - family (default “gaussian”) - response distribution and link function
 - prior - override BRMS defaults
 - brm function

BRMS Description of a Multilevel Model

$$y_i \sim D(f^{-1}(\eta_i), \theta)$$

- `brm` formula argument specifies η , the linear predictor
 - η can be rewritten $X\beta + Zu$ where
 - β and u are population-level and group-level coefficients respectively
- `data` argument specifies design matrices X , Z and the modeled data y .
- `family` argument specifies both
 - D , the distributional family
 - f , the *inverse link function*
- `prior` argument overrides defaults on θ , the family-specific parameter(s)

The Stan Language

- A Stan program
 - Declares data and (constrained) parameter variables
 - Defines log posterior (or penalized likelihood)
 - Computes quantities of interest
- Syntax
 - Influenced by BUGS (plus Java/C++ punctuation)
 - Explicit variable types (like Java/C++, not like Python, R)
 - Named program blocks - distinguish between data and parameters (not like BUGS)
 - Control flow: `if` statements - dynamic branch points (more powerful than BUGS)
- Mathematical operations
 - Stan language includes a very large set of probability distributions and mathematical functions from Stan's math library
 - Efficient, vectorized (almost always)

Stan Example: Laplace's Model of Birth Rate by Sex

Laplace's data on live births in Paris from 1745–1770:

<i>sex</i>	<i>live births</i>
female	241 945
male	251 527

- *Question 1* (Estimation): What is the birth rate of boys vs. girls?
- *Question 2* (Event Probability): Is a boy more likely to be born than a girl?

Laplace computed this analytically. Let's use Stan's NUTS-HMC sampler instead.

Laplace's Model in Stan

```
transformed data {  
  int male = 251527;  
  int female = 241945;  
}  
parameters {  
  real<lower=0, upper=1> theta;  
}  
model {  
  male ~ binomial(male + female, theta);  
}  
generated quantities {  
  int<lower=0, upper=1> theta_gt_half = (theta > 0.5);  
}
```

Laplace's Answer

```
births_model = cmdstan_model("laplace.stan") # compile model
births_fit = births_model$sample()           # run inference algorithm
as.data.frame(births_fit$summary())           # manage results
```

variable	mean	median	sd	q5	q95
theta	0.51	0.51	0.000725	0.509	0.511
theta_gt_half	1.00	1.00	0.000000	1.000	1.000

- *Question 1* (Estimation): What is the birth rate of boys vs. girls?
 θ is 90% certain to lie in (0.509, 0.511)
- *Question 2* (Event Probability): Is a boy more likely to be born than a girl?
Laplace “morally certain” boys more prevalent

Stan Program File

A Stan program consists of one or more named program blocks, strictly ordered

```
functions {  
  // declare, define functions  
} data {  
  // declare input data  
} transformed data {  
  // transform inputs, define program data  
} parameters {  
  // declare (continuous) parameters  
} transformed parameters {  
  // define derived parameters  
} model {  
  // compute the log joint distribution  
} generated quantities {  
  // define quantities of interest  
}
```

Stan Program Blocks - Execution During Sampling

- data, transformed data blocks - executed once on startup
- parameters -
 - on startup: initialize parameters
 - at every step of inference algorithm: validate constraints
- transformed parameters, model blocks - executed every *step* of the sampler
- generated quantities - executed every *iteration* of the sampler
- After every sampler iteration, program outputs current values of all variables in parameters, transformed parameters, and generated quantities blocks.

The Stan Language

- Variables have strong, static types
 - Strong: only values of that type will be assignable to the variable, (type promotion allowed, e.g., `int` to `real`, `complex`)
 - Static: variable type is constant throughout program
 - Types: `vector`, `row_vector`, `matrix`, `complex`, `array`, and `tuple`
 - Variables can be declared with constraints on `upper` and/or `lower` bounds
- Motivation
 - Static types make programs easier to comprehend, debug, and maintain
 - Programming errors can be detected at compile-time
 - Constrained types catch runtime errors

```
int<lower=0> N, N_time;           // constraint applies to both N, N_time
vector[3]<lower=0> sigma;
array[3] matrix[M, N] some_xs;
array[N] int<lower=0, upper=N_time> time; // use arrays for structured int data
```

The Stan Language

The `model` block

- defines the joint log probability density function given parameters
- this function is the sum of all distribution and log probability increment statements in the model block

Distribution statements

```
y ~ normal(mu, sigma);  
mu ~ normal(0, 10);  
sigma ~ normal(0, 1);
```

Log probability increment statements

```
target += normal_lupdf(y | mu, sigma);  
target += normal_lupdf(mu | 0, 10);  
target += normal_lupdf(sigma | 0, 1);
```


Processing Steps

- `births_model = cmdstan_model("laplace.stan")` `# compile model`
 - Stan compiler translates Stan file to C++ file
 - C++ file is compiled to executable program, via GNU Make
- `births_fit = births_model$sample()` `# do inference`
 - Interfaces run the compiled executable (as a subprocess) and manage the per-chain outputs (modified CSV format files)
- `as.data.frame(births_fit$summary())` `# manage results`
 - Parse CSV outputs into in-memory object
 - Access individual parameters and quantities of interest
 - Run summary and diagnostic reports

From BRMS to Stan

- In BRMS, specify arguments to the `brm` function: formula, data, family, prior.
BRMS supplies defaults for all elements of the model

$$y_i \sim D(f^{-1}(\eta_i), \theta)$$

- D, f^{-1}, θ are the distributional family, link function, and distributional parameters
 - η is the regression formula
- In Stan, specify a model using elements of the Stan probabilistic programming language.
 - `data` block defines y , as well as all unmodeled data inputs
 - `parameters` block defines all distributional parameters.
 - `model` block specifies the likelihood and priors. (formula, distributional family, priors).

Model Development: Hello, World!

- A "Hello, World!" program is the name given to the first, simplest possible program written when learning a new programming language.
- Our “Hello, World!” model is the **complete pooling model**
all subjects are alike; reaction time is predicted by days since study start.
- BRMS formula: `Reaction ~ Days`
- Stan likelihood: `y ~ normal(alpha + days_c * beta, sigma);`
 - center predictors
 - use weakly informative priors

Stan Model

```
data {  
  int<lower=0> N;   vector[N] days;   vector[N] y;  // reaction time  
}  
  
transformed data { // center predictor variable "days"  
  real days_mean = mean(days); vector[N] days_c = days - days_mean;  
}  
  
parameters {  
  real alpha;  real beta;  // intercept, slope  
  real<lower=0> sigma; // residual standard deviation  
}  
  
model {  
  y ~ normal(alpha + days_c * beta, sigma);  
  alpha ~ normal(250, 50); // informed prior for human reaction times in ms  
  beta ~ normal(10, 10); // weakly informed prior for per-day effect  
  sigma ~ normal(0, 10); // very weakly informative prior  
}  
  
generated quantities { // adjust for centered predictor values  
  real intercept = alpha - beta * days_mean;  
}
```

Reading for Next Time

- Gelman blogpost [varying-slope/varying-intercept models](#)
“If you want to program this directly in Stan, once you have varying intercepts and slopes, you have to deal with covariance-matrix decompositions and arrays of coefficient vectors, and it’s all a hairy mess.”
- Bob Carpenter’s response [Varying slopes and intercepts in Stan: still painful in 2024](#)
“I have to confess up front that Andrew’s right—this is still painful”
- [Optimization through Cholesky Factorization](#) Stan User’s Guide, Regression Models
- [QR Decomposition](#) Michael Betancourt case study