

From BRMS to Stan

Mitzi Morris

2025-03-25

Stan Development Team



BRMS: Bayesian Regression and Multilevel modeling in Stan

Recap from last week * BRMS extended formula syntax for multi-level regressions

`response ~ pterms + (gterms | group)`

- *pterms* define *population-level* effects; same for all observations.
- *gterms* define *group-level* effects; vary across *group* variable.
- The intercept term is 1 or 0 for no intercept; if unspecified, default is 1.

```
fit1 <- brm(Reaction ~ Days + (Days | Subject), data = sleepstudy)
```

equivalent to

```
fit1 <- brm(Reaction ~ 1 + Days + (1 + Days | Subject), data = sleepstudy)
```

Stan Program File (review)

A Stan program consists of one or more named program blocks, strictly ordered

```
functions {  
  // declare, define functions  
} data {  
  // declare input data  
} transformed data {  
  // transform inputs, define program data  
} parameters {  
  // declare (continuous) parameters  
} transformed parameters {  
  // define derived parameters  
} model {  
  // compute the log joint distribution  
} generated quantities {  
  // define quantities of interest  
}
```

Stan Program Blocks - Execution During Sampling (review)

- data, transformed data blocks - executed once on startup
- parameters -
 - on startup: initialize parameters
 - at every step of inference algorithm: validate constraints
- transformed parameters, model blocks - executed every *step* of the sampler
- generated quantities - executed every *iteration* of the sampler
- After every sampler iteration, program outputs current values of all variables in parameters, transformed parameters, and generated quantities blocks.

From BRMS to Stan

- BRMS: specify arguments to the `brm` function: formula, data, family, prior.
 - BRMS generates Stan model code.
 - BRMS code is not quite human-readable, but efficiently coded.
- *Goal: write efficient, robust Stan program*
 - map regression formula to Stan program's sampling distribution statement.
 - data block defines all data inputs - outcomes and predictors, plus dimensions.
 - transformed data block mean-centers predictor variables.
 - parameters block defines all distributional parameters.
 - model block specifies the likelihood and priors.
- When should you do this?
 - When model specification in BRMS is long / complicated / not quite possible.

Stepwise Model Development: Hello, World!

- A "Hello, World!" program is the name given to the first, simplest possible program written when learning a new programming language.
 - **Pro tip: always start with “Hello, World!”**
- End goal is a efficient and maintainable multi-level model
 $\text{Reaction} \sim \text{Days} + (\text{Days} | \text{Subject})$
- Initial goal is a simple linear model - complete pooling across subjects
 $\text{Reaction} \sim \text{Days}$
($\text{Reaction} \sim 1 + \text{Days}$ - by default, model includes global intercept.)
- Carry over BRMS efficiencies to Stan model

Stan Model sleep_simple.stan

```
data {  
  int<lower=0> N;    vector[N] day;    vector[N] y;  // reaction time  
}  
transformed data {  
  real day_mean = mean(day);  
  vector[N] day_centered = day - day_mean;  
}  
parameters {  
  real alpha;  real b_day;  // intercept, slope  
  real<lower=0> sigma;  // residual standard deviation  
}  
model {  
  y ~ normal(alpha + day_centered * b_day, sigma);  
  alpha ~ normal(250, 50);  // informed prior for human reaction times in ms  
  b_day ~ normal(10, 10);  // weakly informed prior for per-day effect  
  sigma ~ normal(0, 10);  // very weakly informative prior  
}  
generated quantities {  
  real b_intercept = alpha - b_day * day_mean;  
  vector[N] y_rep;  vector[N] log_lik;  // for posterior predictive checks  
  // ...  
}
```

Notebook - Stan, BRMS complete pooling model

```
sleep_data = list(  
  N = nrow(sleepstudy), J = length(unique(sleepstudy$Subject)),  
  subj = as.integer(sleepstudy$Subject), day = sleepstudy$Days,  
  y = as.double(sleepstudy$Reaction)  
)  
  
sleep_simple = cmdstan_model("stan/sleep_simple.stan")  
sleep_simple_stanfit = sleep_simple$sample(data = sleep_data)  
as.data.frame(sleep_simple_stanfit$summary(variables = c('b_intercept', 'b_day', 'sigma')))  
  
priors <- c(set_prior("normal(250, 50)", class = "Intercept"),  
  set_prior("normal(10, 10)", class = "b"),  
  set_prior("normal(0, 10)", class = "sigma"))  
  
sleep_simple_brmsfit = brm(Reaction ~ Days, data = sleepstudy, prior = priors)  
sleep_simple_brmsfit
```


Multilevel models

- Specify model in terms of the inherent structure of the data
- Sleep study: reaction time varies by subject
 - BRMS formula: `Reaction ~ 1 + Days + (1 + Days|Subject)`
- Expand Stan model:
 - Predictor vector β is *multivariate normal*
 - `b_subj ~ multi_normal(mu_subj, sigma_subj)`
 - `mu_subj` is vector, `sigma_subj` is covariance matrix.
- *Problems*
 - Stan distribution `multi_normal` - requires inverting covariance matrix at every evaluation - computationally expensive.
 - Two sources of variance: `sigma` and hierarchical variance `sigma_subj`
difficult to estimate from small number of observations per group.

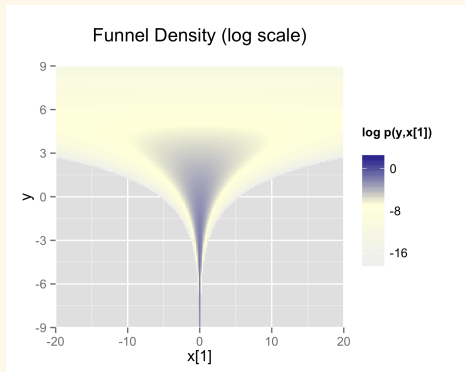
Multilevel models

- **Partial Pooling:** The hierarchical prior controls the pooling between levels:
 - Similar data across levels \rightarrow Low hierarchical variance \rightarrow Strong pooling
 - Dissimilar data across levels \rightarrow High hierarchical variance \rightarrow Weak pooling
- **Problem:** For low numbers of observations, MCMC sampler cannot resolve residual variance `sigma` and variance of hierarchical prior `sigma_subj`; many divergences
- **Solution:** Reparameterization, following code example in Stan User's Guide section [Hierarchical models and the non-centered parameterization](#)

Multilevel models and Neal's Funnel

- Neal's Funnel: extreme example of a challenging hierarchical prior

$$p(y, x) = \text{normal}(y \mid 0, 3) \times \prod_{n=1}^9 \text{normal}(x_n \mid 0, \exp(y/2)).$$



To explore neck of the funnel:

- small steps on x-axis, large steps on y-axis.

To explore mouth of the funnel:

- large steps on x-axis, small steps on y-axis.

But stepsize is same for all axes;
cannot adequately sample either.

Funnel Example: Stan Implementations

Centered parameterization

- “Natural” parameterization.

```
parameters {  
  real y;  
  vector[9] x;  
}  
  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```

Non-centered parameterization


- Parameters block: declare standardized parameters.
- Transformed parameters: declare *variables* add offset (location), multiply by scale.

```
parameters {  
  real y_raw;  
  vector[9] x_raw;  
}  
  
transformed parameters {  
  // offset is 0, just multiply by scale  
  real y = 3.0 * y_raw;  
  vector[9] x = exp(y/2) * x_raw;  
}  
  
model {  
  y_raw ~ std_normal(); // y ~ normal(0, 3)  
  x_raw ~ std_normal(); // x ~ normal(0, exp(y/2))  
}
```

Compare Funnel Fits


Centered parameterization

- Cannot explore neck of funnel
many divergences (in orange)



Non-centered parameterization

- Explores further; no divergences



BRMS to Stan

- Recap: formula `Reaction ~ 1 + Days` becomes distribution statement:
`y ~ normal(alpha + day_centered * b_day, sigma);`
- Coding `Reaction ~ 1 + Days + (1 + Days | Subject)`
 - Global intercept and slope for Days
 - Subject-specific random intercepts and slopes
- Correlation between random effects: prior on subject effect is a multivariate normal with mean vector μ and covariance matrix Σ .
- Multivariate reparameterization

Multi-variate reparameterization

Following blogpost [Varying slopes and intercepts in Stan: still painful in 2024](#)

Centered parameterization

```
parameters {  
  vector[K] beta;  
  vector[K] mu;  
  cov_matrix[K] Sigma;  
  // ...  
}  
  
model {  
  beta ~ multi_normal(mu, Sigma);  
  // ...  
}
```

Non-centered parameterization

```
parameters {  
  vector<lower=0>[K] tau;  
  cholesky_factor_corr[K] L_Omega;  
  matrix[K, J] beta_std;  
}  
  
transformed parameters {  
  matrix[J, K] beta =  
    (diag_pre_multiply(tau, L_Omega) * beta_std)';  
}  
  
model {  
  // non-centered priors  
  tau ~ exponential(1);  
  L_Omega ~ lkj_corr_cholesky(K);  
  to_vector(beta_std) ~ std_normal();  
}
```

Specifying the Likelihood

- $\text{Reaction} \sim 1 + \text{Day} + (1 + \text{Subject} \mid \text{Day})$
- Create design matrix x with column 1 for group-level intercept term

```
transformed data {  
  matrix[N, 2] x;  
  x[ , 1] = rep_vector(1, N);  
  x[ , 2] = day;  
}
```

- Add group-level term to regression formula.

```
vector[N] eta = b_intercept + b_day * day + rows_dot_product(x, beta[ subj, ]);  
y ~ normal(eta, sigma);
```


Parameters, transformed parameters, model blocks

```
parameters {  
  real b_intercept; real b_day; real<lower=0> sigma;  
  
  vector<lower=0>[2] tau; cholesky_factor_corr[2] L_Omega;  
  matrix[2, J] beta_std;  
}  
transformed parameters {  
  // random effects matrix scaled, transposed (centered at 0)  
  matrix[J, 2] beta = (diag_pre_multiply(tau, L_Omega) * beta_std)';  
}  
model {  
  vector[N] eta = b_intercept + b_day * day + rows_dot_product(x, beta[ subj, ]);  
  y ~ normal(eta, sigma);  
  
  b_intercept ~ normal(250, 50); b_day ~ normal(10, 10); sigma ~ exponential(1);  
  
  tau ~ exponential(1); L_Omega ~ lkj_corr_cholesky(2);  
  to_vector(beta_std) ~ std_normal();  
}
```

Recovering the Quantities of Interest

```
generated quantities {  
  // match BRMS outputs  
  real sd_intercept = tau[1];  
  real sd_day = tau[2];  
  
  // Reconstruct correlation matrix  
  matrix[2, 2] Omega;  
  Omega = multiply_lower_tri_self_transpose(L_Omega);  
  real cor_intercept_day = Omega[1, 2];  
  
  // Posterior likelihood and posterior predictive y-replicates  
  vector[N] y_rep;  vector[N] log_lik;  
  { // don't save to output  
    vector[N] eta = b_intercept + b_day * day + rows_dot_product(x, beta[ subj, ] );  
    y_rep = to_vector(normal_rng(eta, sigma));  
    for (n in 1:N) {  
      log_lik[n] = normal_lpdf(y[n] | eta[n], sigma);  
    }  
  }  
}
```

Notebook Demo

```
# Stan multilevel model
```

```
sleep_mlm = cmdstan_model(stan_file = "stan/sleep_mlm.stan")
sleep_mlm_stanfit = sleep_mlm$sample(data = sleep_data, ... )
as.data.frame(sleep_mlm_stanfit$summary(
  variables = c('b_intercept', 'b_day', 'sigma',
    'sd_intercept', 'sd_day', 'cor_intercept_day'))))
```

```
# BRMS multilevel model
```

```
priors <- c(
  set_prior("normal(250, 50)", class = "Intercept"),
  set_prior("normal(10, 10)", class = "b"),
  set_prior("exponential(1)", class = "sigma"),
  set_prior("exponential(1)", class = "sd"),
  set_prior("lkj_corr_cholesky(2)", class = "cor"),
)
sleep_mlm_brmsfit <- brm(Reaction ~ Days + (Days|Subject),
  data = sleepstudy,
  prior = priors)
sleep_mlm_brmsfit
```

Discussion

- BRMS formula syntax provides concise description of the regression.
 - function `brm` generates Stan code given both the formula and the data
 - default is a simple linear model.
- BRMS imposes choices; when overriding them becomes too complicated
just code it in Stan!
- An efficient Stan program makes it easy for the sampler to explore the posterior density.
 - zero-center predictors, make sure they are on the same scale.
 - the choice of the centered vs. non-centered parameterization depends on the amount of observations per group-level predictor.
 - use the non-centered parameterization for low-data regimes.

References

Stan User's Guide:

- Efficiency Tuning, [Hierarchical models and the non-centered parameterization](#)
- Efficiency Tuning, [Multivariate reparameterization](#)
- Regression, [Multivariate regression example](#)

Many Thanks!

Questions???