

From BRMS to Stan

Mitzi Morris

2025-05-07

Stan Development Team



My Questions for You

- What kinds of models do you use for research? For teaching?
- BRMS
 - How often do you use it: never, sometimes, always?
 - What do you like/dislike/find confusing about BRMS?
- Stan
 - Have you run Stan models through an interface, RStan, CmdStanR, CmdStanPy?
 - Have you written / tried to write or modify Stan models?
 - What do you like/dislike/find confusing about Stan?

Talk Outline

- Brief discussion of BRMS
- Stan from 0 to 60 in 30 minutes
 - Model 1: “Hello World”
 - Model 2: black diamond - varying-slope / varying intercept
- Notebook: github.com/mitzimorris/brms2stan/brms2stan.qmd

BRMS: Bayesian Regression and Multilevelmodeling in Stan



brms

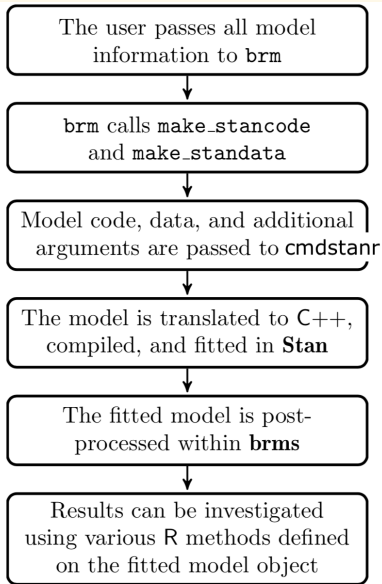
Bayesian regression models using Stan

- Fit Bayesian generalized linear / non-linear multivariate multilevel models using Stan.
- Specify model via lme4-like formula syntax.

```
sleep_mlm_brmsfit <-  
  brm(Reaction ~ Days + (Days|Subject),  
      data = sleepstudy)
```

<https://paul-buerkner.github.io/brms/>

BRMS Processing



BRMS Processing

- Pass all information to BRMS

```
brm(Reaction ~ Days + (Days|Subject), data = sleepstudy)
```

- Get results

```
Family: gaussian
```

```
Links: mu = identity; sigma = identity
```

```
Formula: Reaction ~ Days + (Days | Subject)
```

```
Data: sleepstudy (Number of observations: 180)
```

Multilevel Hyperparameters:

~Subject (Number of levels: 18)

	Estimate	Est.Error
sd(Intercept)	27.27	7.03
sd(Days)	6.60	1.51
cor(Intercept,Days)	0.07	0.30

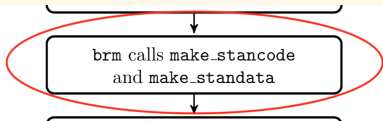
Regression Coefficients:

	Estimate	Est.Error
Intercept	251.31	7.40
Days	10.45	1.70
sigma	25.87	1.52

Reasons to Use BRMS

- Familiar syntax (for R users).
- Handles many common models: IRT, zero-inflated, additive distributional, phylogenetic, missing data imputation.
 - See [Vignettes and Articles](#)
- Sound and robust implementations, e.g.
 - Zero-centers predictors (like RStanARM).
 - Makes varying slope/ varying intercept models easy.
- Large community of users
 - Majority of questions on Stan mailing list are BRMS.
 - Many helpful community experts to answer them!

BRMS Drawbacks



- Generated Stan code is hard to read/modify.
- Priors used may be suboptimal.

See [github.com/stan-dev/stan](https://github.com/stan-dev/stan/wiki/prior-choice-recommendations) wiki page [Prior Choice Recommendations](https://github.com/stan-dev/stan/wiki/prior-choice-recommendations)

A screenshot of a web browser showing the GitHub page for "Prior Choice Recommendations" under the "stan-dev / stan" repository. The browser's address bar shows the URL "github.com/stan-dev/stan/wiki/prior-choice-recommendations". The GitHub navigation bar includes links for "Code", "Issues" (133), "Pull requests" (12), "Discussions", "Actions", "Projects", "Wiki" (which is underlined), "Security", and "Insights". A search bar on the right says "Type / to search". The main heading of the page is "Prior Choice Recommendations", followed by the text "Aki Vehtari edited this page last week · [58 revisions](#)". Below this, the section "5 levels of priors" is visible, with the first item being "Flat prior (not usually recommended):".

BRMS Priors

Prior Definitions for brms Models

```
> prior_summary(sleep_mlm_brmsfit)
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub	source
	(flat)	b								default
	(flat)	b	Days							(vectorized)
student_t(3, 288.7, 59.3)		Intercept								default
lkj_corr_cholesky(1)		L								default
lkj_corr_cholesky(1)		L		Subject						(vectorized)
student_t(3, 0, 59.3)		sd						0		default
student_t(3, 0, 59.3)		sd		Subject				0		(vectorized)
student_t(3, 0, 59.3)		sd	Days	Subject				0		(vectorized)
student_t(3, 0, 59.3)		sd	Intercept	Subject				0		(vectorized)
student_t(3, 0, 59.3)		sigma						0		default

- To change the priors you must understand the model structure.
- If you understand the model structure, you can code it in Stan!

Changing BRMS Priors

```
priors <- c(
  set_prior("normal(250, 50)", class = "Intercept"),
  set_prior("normal(10, 10)", class = "b"),
  set_prior("exponential(1)", class = "sigma"),
  set_prior("exponential(1)", class = "sd"),
  set_prior("lkj_corr_cholesky(2)", class = "cor"),
)
sleep_mlm_brmsfit <- brm(Reaction ~ Days + (Days|Subject),
  data = sleepstudy,
  prior = priors)
```

Easy to get wrong - cf. questions on discourse.mc-stan.org, e.g.: recent discussion

Equally complicated - custom families, custom link functions.

From BRMS to Stan

When model specification in BRMS is long / complicated / not quite possible - try Stan.

- How do you write a Stan program?
- BRMS Formula provides structure of regression and data inputs.
- Mapping BRMS elements to a Stan program
 - `data` block defines all data inputs - outcomes and predictors, plus dimensions.
 - `parameters` block defines all distributional parameters.
 - `model` block specifies the likelihood and priors.

Stan Program Structure

A Stan program consists of one or more named program blocks, strictly ordered.

```
functions {  
  // declare, define functions  
} data {  
  // declare input data  
} transformed data {  
  // transform inputs, define program data  
} parameters {  
  // declare (continuous) parameters  
} transformed parameters {  
  // define derived parameters  
} model {  
  // compute the log joint distribution  
} generated quantities {  
  // define quantities of interest  
}
```

Stan Program Execution

- NUTS-HMC sampler discretizes Hamiltonian dynamics into steps.
 - 1 sampler iteration is comprised of many steps.
 - Good NUTS-HMC sampler intro: [Faster estimation of Bayesian models in ecology using Hamiltonian Monte Carlo](#)
- data, transformed data blocks - executed once on startup - *cheap*
- parameters -
 - on startup: initialize parameters
 - at every step of inference algorithm: validate constraints
- transformed parameters, model blocks - executed every *step* of the sampler - *expensive*
- generated quantities - executed every *iteration* of the sampler - *less expensive*

Stepwise Model Development: Hello, World!

- A "Hello, World!" program is the name given to the first, simplest possible program written when learning a new programming language.
 - **Pro tip: always start with “Hello, World!”**
- End goal is a efficient and maintainable multi-level model
 - `Reaction ~ Days + (Days|Subject)`
- Initial goal is a simple linear model - ignore difference between subjects
 - BRMS formula: `Reaction ~ Days`
- Simple linear regression: $y \sim \alpha + \beta * x$,
 - where reaction time is y and day is x .

Initial Stan Model sleep_simple.stan

```
data {  
  int<lower=0> N;  
  vector[N] day;  
  vector[N] y;  // reaction time  
}  
parameters {  
  real alpha;  // intercept  
  real b_day;  // slope  
  real<lower=0> sigma; // residual standard deviation  
}  
model {  
  y ~ normal(alpha + day * b_day, sigma);  
  alpha ~ normal(250, 50);  // informed prior for human reaction times in ms  
  b_day ~ normal(10, 10);  // weakly informed prior for per-day effect  
  sigma ~ normal(0, 10);  // very weakly informative prior  
}
```

Improved Stan Model sleep_simple.stan - Mean-center Data

This decreases the correlation between the intercept and slope.
(BRMS, RStanArm do this automatically.)

```
data {  
  int<lower=0> N;   vector[N] day;   vector[N] y;  
}  
transformed data {  
  real day_mean = mean(day);  
  vector[N] day_centered = day - day_mean;    // mean-center data  
}  
parameters {  
  real alpha;  real b_day;  real<lower=0> sigma;  
}  
model {  
  y ~ normal(alpha + day_centered * b_day, sigma);  
  alpha ~ normal(250, 50); b_day ~ normal(10, 10); sigma ~ normal(0, 10);  
}  
generated quantities {  
  real b_intercept = alpha - b_day * day_mean;    // recover intercept  
}
```


BRMS stancode for Reaction ~ Days

```
> stancode(sleep_simple_brmsfit)
// generated with brms 2.22.0
functions {
}
data {
  int<lower=1> N; // total number of observations
  vector[N] Y; // response variable
  int<lower=1> K; // number of population-level effects
  matrix[N, K] X; // population-level design matrix
  int<lower=1> Kc; // number of population-level effects
  int prior_only; // should the likelihood be ignored?
}
transformed data {
  matrix[N, Kc] Xc; // centered version of X without an intercept
  vector[Kc] means_X; // column means of X before centering
  for (i in 2:K) {
    means_X[i - 1] = mean(X[, i]);
    Xc[, i - 1] = X[, i] - means_X[i - 1];
  }
}
```

```
parameters {
  vector[Kc] b; // regression coefficients
  real Intercept; // temporary intercept for centered predictor
  real<lower=0> sigma; // dispersion parameter
}
transformed parameters {
  real lprior = 0; // prior contributions to the log posterior
  lprior += normal_lpdf(b | 10, 10);
  lprior += normal_lpdf(Intercept | 250, 50);
  lprior += normal_lpdf(sigma | 0, 10)
    - 1 * normal_lccdf(0 | 0, 10);
}
model {
  if (!prior_only) { // likelihood including constants
    target += normal_id_glm_lpdf(Y | Xc, Intercept, b, sigma);
  }
  target += lprior; // priors including constants
}
generated quantities { // actual population-level intercepts
  real b_Intercept = Intercept - dot_product(means_X, b);
}
```

Multilevel Models

- Account for the structure in the data.
 - Individual observations are drawn from one or more common populations.
 - Estimate population-level mean, variance as well as individual-level mean, variance.
 - Population level parameters inform priors for individual level parameters.
- Provide **partial pooling** of information.
 - The hierarchical prior controls the pooling between levels.
 - Similar data across levels → Low hierarchical variance → strong pooling.
When hierarchical variance approaches 0, complete pooling.
 - Dissimilar data across levels → High hierarchical variance → weak pooling.
When hierarchical variance approaches ∞ , no pooling.


Multilevel models

- Expand Stan model:
 - A per-subject parameter vector β (`b_subj`) is *multivariate normal*
 - `b_subj ~ multi_normal(mu_subj, sigma_subj)`
 - `mu_subj` is vector, of length number of subjects,
`sigma_subj` is covariance matrix.
- *Problems*
 - Stan distribution `multi_normal` - requires inverting covariance matrix at every evaluation - computationally expensive.
 - Two sources of variance: `sigma` and hierarchical variance `sigma_subj`
difficult to estimate from small number of observations per group.
Without sufficient data → **funnel distribution**
- **Solution:** Reparameterization, following code example in Stan User's Guide section *Hierarchical models and the non-centered parameterization*

Multilevel models and the Funnel

Neal's Funnel: extreme example of a challenging hierarchical prior

$$p(y, x) = \text{normal}(y \mid 0, 3) \times \prod_{n=1}^9 \text{normal}(x_n \mid 0, \exp(y/2)).$$



To explore neck of the funnel:

- small steps on x -axis, large steps on y -axis.

To explore mouth of the funnel:

- large steps on x -axis, small steps on y -axis.

But stepsize is same for all axes;
cannot adequately sample either.

Prior distribution; no data.

Funnel Example: Stan Implementations

Centered parameterization

- “Natural” parameterization.

```
parameters {  
  real y;  
  vector[9] x;  
}  
  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```

Note: minimal model, no predictors

- variable y == hierarchical variance
- variable x == group-level covariate

Non-centered parameterization

- Parameters block: declare standardized parameters.
- Transformed parameters: declare *variables* add offset (location), multiply by scale.

```
parameters {  
  real y_raw;  
  vector[9] x_raw;  
}  
  
transformed parameters {  
  // offset is 0, just multiply by scale  
  real y = 3.0 * y_raw;  
  vector[9] x = exp(y/2) * x_raw;  
}  
  
model {  
  y_raw ~ std_normal(); // y ~ normal(0, 3)  
  x_raw ~ std_normal(); // x ~ normal(0, exp(y/2))  
}
```

Stan's Affine Transform

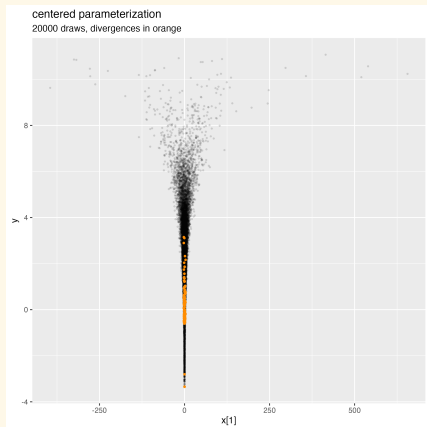
Stan's [Affinely Transformed Scalar](#) can be applied element-wise to vector x to facilitate the [non-centered parameterization](#)

```
parameters {  
  real<multiplier = 3.0> y;  
  vector<multiplier = exp(y/2)>[9] x;  
}  
  
model {  
  y ~ std_normal(); //  $y \sim \text{normal}(0, 3)$   
  x ~ std_normal(); //  $x \sim \text{normal}(0, \exp(y/2))$   
}
```

Compare Funnel Fits

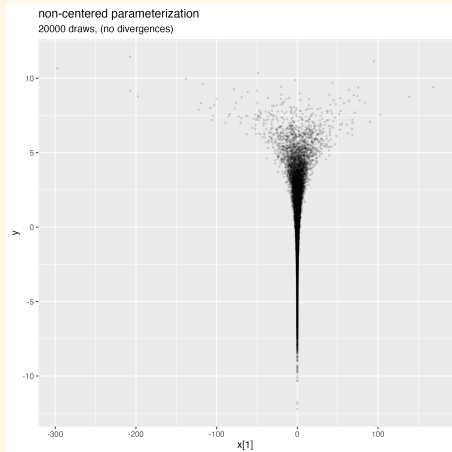
Centered parameterization

- Cannot explore neck of funnel
many divergences (in orange)



Non-centered parameterization

- Explores further; no divergences



BRMS to Stan

- “Hello, World!” model - formula `Reaction ~ 1 + Days`
 - coded as distribution statement:
`y ~ normal(alpha + day_centered * b_day, sigma);`
- Target model - formula `Reaction ~ 1 + Days + (1 + Days | Subject)`
 - Global intercept and slope for Days
 - Subject-specific random intercepts and slopes
- Correlation between random effects: prior on subject effect is a multivariate normal with mean vector μ (mu) and covariance matrix Σ (Sigma).
- See Gelman & Hill 2007, chapter 11, 12, 13 or [this blogpost](#)
- [Multivariate reparameterization](#)

Specifying the Likelihood

- $\text{Reaction} \sim 1 + \text{Day} + (1 + \text{Subject} \mid \text{Day})$
- Create design matrix x with column 1 for group-level intercept term

```
transformed data {  
  matrix[N, 2] x;  
  x[ , 1] = rep_vector(1, N);  
  x[ , 2] = day;  
}
```

- Add group-level term to regression formula.

```
vector[N] eta = b_intercept + b_day * day + rows_dot_product(x, beta[ subj, ]);  
y ~ normal(eta, sigma);
```

Multi-variate reparameterization

Centered parameterization

```
data {  
  int K; // num predictors + intercept  
  int J; // num groups  
}  
  
parameters {  
  vector[K, J] beta; // group covariates  
  vector[K] mu; // location  
  cov_matrix[K] Sigma; // scale  
  // ...  
}  
  
model {  
  // hierarchical prior  
  for (j in 1:J) {  
    beta[, j] ~ multi_normal(mu, Sigma);  
  }  
  // ...  
}
```

Non-centered parameterization

```
data {  
  int K; int J;  
}  
  
parameters {  
  vector<lower=0>[K] tau;  
  cholesky_factor_corr[K] L_Omega;  
  matrix[K, J] beta_std;  
}  
  
transformed parameters {  
  matrix[J, K] beta =  
    (diag_pre_multiply(tau, L_Omega) * beta_std)';  
}  
  
model {  
  // non-centered priors  
  tau ~ exponential(1);  
  L_Omega ~ lkj_corr_cholesky(K);  
  to_vector(beta_std) ~ std_normal();  
}
```

Parameters, transformed parameters, model blocks

```
parameters {  
  real b_intercept; real b_day; real<lower=0> sigma;  
  
  vector<lower=0>[2] tau; cholesky_factor_corr[2] L_Omega;  
  matrix[2, J] beta_std;  
}  
transformed parameters {  
  // random effects matrix scaled, transposed (centered at 0)  
  matrix[J, 2] beta = (diag_pre_multiply(tau, L_Omega) * beta_std)';  
}  
model {  
  vector[N] eta = b_intercept + b_day * day + rows_dot_product(x, beta[ subj, ]);  
  y ~ normal(eta, sigma);  
  
  b_intercept ~ normal(250, 50); b_day ~ normal(10, 10); sigma ~ exponential(1);  
  
  tau ~ exponential(1); L_Omega ~ lkj_corr_cholesky(2);  
  to_vector(beta_std) ~ std_normal();  
}
```

Recovering the Quantities of Interest

```
generated quantities {  
  // match BRMS outputs  
  real sd_intercept = tau[1];  
  real sd_day = tau[2];  
  
  // Reconstruct correlation matrix  
  matrix[2, 2] Omega;  
  Omega = multiply_lower_tri_self_transpose(L_Omega);  
  real cor_intercept_day = Omega[1, 2];  
  
  // Posterior likelihood and posterior predictive y-replicates  
  vector[N] y_rep;  vector[N] log_lik;  
  { // don't save to output  
    vector[N] eta = b_intercept + b_day * day + rows_dot_product(x, beta[ subj, ] );  
    y_rep = to_vector(normal_rng(eta, sigma));  
    for (n in 1:N) {  
      log_lik[n] = normal_lpdf(y[n] | eta[n], sigma);  
    }  
  }  
}
```

Notebook

- github.com/mitzimirris/brms2stan/brms2stan.qmd
- Let's run this in RStudio

References

Stan User's Guide:

- Efficiency Tuning, [Hierarchical models and the non-centered parameterization](#)
- Efficiency Tuning, [Multivariate reparameterization](#)
- Regression, [Multivariate regression example](#)
- Regression, [Optimization through Cholesky Factorization](#)

Andrew Gelman's Blog

- Gelman blogpost [varying-slope/varying-intercept models](#)
- Bob Carpenter's response [Varying slopes and intercepts in Stan: still painful in 2024](#)

Many Thanks!

Questions???